

Need for DAOS

- Object storage for unstructured data
- Complex I/O patterns
- Vast quantities of random reads and writes
- Critical data access time

Objectives

1. Installation of DAOS Server and Client RPMs
2. Orchestrate Object Storage
3. Demonstrate User Interactions

Objective 1: Installation of DAOS Server and Client RPMs

Setting Network Adaptor

- Two adaptors used:
 - For SSH connection: used host-only adapter
 - For VM to VM connection: used NAT network adapter

```
[root@server ~]# ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether 08:00:27:c1:40:b8 brd ff:ff:ff:ff:ff:ff
    inet 192.168.100.9/24 brd 192.168.100.255 scope global noprefixroute dynamic enp0s3
        valid_lft 594sec preferred_lft 594sec
    inet6 fe80::ee83:e7ce:45ca:f9b/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether 08:00:27:b4:3a:98 brd ff:ff:ff:ff:ff:ff
    inet 192.168.56.105/24 brd 192.168.56.255 scope global noprefixroute dynamic enp0s8
        valid_lft 562sec preferred_lft 562sec
    inet6 fe80::8a74:f52f:3e83:e609/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
```

Server installation

- Perform a network scan to find an available provider and its interface

```
[root@server ~]# daos_server network scan
DAOS Server config loaded from /etc/daos/daos_server.yml
-----
localhost
-----

-----
NUMA Socket 0
-----

Provider      Interfaces
-----
ofi+sockets  enp0s3, enp0s8
```

- Update the daos_server.yml file

```
## DAOS server configuration file.
#
name: daos_server
#
access_points: ['192.168.100.9']
port: 10001
#
transport_config:
  allow_insecure: true
  client_cert_dir: /etc/daos/certs/clients
  ca_cert: /etc/daos/certs/daosCA.crt
  cert: /etc/daos/certs/server.crt
  key: /etc/daos/certs/server.key
provider: ofi+sockets
socket_dir: /var/run/daos_server
nr_hugepages: 4096
control_log_mask: DEBUG
control_log_file: /tmp/daos_server.log
helper_log_file: /tmp/daos_admin.log
engines:
-
  targets: 8
  nr_xs_helpers: 0
  fabric_iface: enp0s3
  fabric_iface_port: 31316
  log_mask: INFO
  log_file: /tmp/daos_engine_0.log
  env_vars:
    - CRT_TIMEOUT=30
  storage:
  -
    scm_mount: /mnt/daos0
    class: ram
    scm_size: 16
```

Client installation

- For client: Update the daos_agent.yml file

```
# DAOS agent configuration file.
#
name: daos_server
access_points: ['192.168.100.9']
port: 10001
transport_config:
  allow_insecure: true
  ca_cert: /etc/daos/certs/daosCA.crt
  cert: /etc/daos/certs/agent.crt
  key: /etc/daos/certs/agent.key
log_file: /tmp/daos_agent.log
```

- For admin: Update the daos_control.yml file

```
# DAOS manager (dmg) configuration file.
#
name: daos_server
port: 10001
hostlist: ['192.168.100.9']
transport_config:
  allow_insecure: true
  ca_cert: /etc/daos/certs/daosCA.crt
  cert: /etc/daos/certs/admin.crt
  key: /etc/daos/certs/admin.key
#
```

Starting Server

```
[root@server ~]# systemctl daemon-reload
[root@server ~]# systemctl start daos_server
[root@server ~]# systemctl status daos_server
■ daos_server.service - DAOS Server
   Loaded: loaded (/usr/lib/systemd/system/daos_server.service; enabled; vendor preset: disabled)
   Active: active (running) since Tue 2024-06-11 02:03:30 IST; 1 weeks 5 days ago
 Main PID: 1286 (daos_server)
   CGroup: /system.slice/daos_server.service
           └─1286 /usr/bin/daos_server start
```

Starting Agent

```
[root@client1 ~]# systemctl start daos_agent
[root@client1 ~]# systemctl status daos_agent
■ daos_agent.service - DAOS Agent
   Loaded: loaded (/usr/lib/systemd/system/daos_agent.service; enabled; vendor preset: disabled)
   Active: active (running) since Tue 2024-06-11 00:44:29 IST; 6 days ago
 Main PID: 1277 (daos_agent)
    CGroup: /system.slice/daos_agent.service
            └─1277 /usr/bin/daos_agent

Jun 11 00:44:29 client1 systemd[1]: Started DAOS Agent.
Jun 11 00:44:29 client1 daos_agent[1277]: DAOS Agent v2.0.3 (pid 1277) listening on /var/run/d...ock
Hint: Some lines were ellipsized, use -l to show in full.
[root@client1 ~]#
```


Storage format and system query

```
[root@client1 ~]# dmfg storage format -l 192.168.100.9 --force -i
```

Format Summary:

Hosts	SCM Devices	NUMe Devices
-------	-------------	--------------

192.168.100.9	1	0
---------------	---	---

```
[root@client1 ~]# dmfg system query -v
```

Rank	UUID	Control Address	Fault Domain	State	Reason
0	c3f36181-ec0a-4976-9912-3354fd6e56ca	192.168.100.9:10001	/server	Joined	

```
[root@client1 ~]# dmfg storage query usage
```

Hosts	SCM-Total	SCM-Free	SCM-Used	NUMe-Total	NUMe-Free	NUMe-Used
192.168.100.9	17 GB	17 GB	0 %	0 B	0 B	N/A

Objective 2

Orchestrate Object Storage

Orchestrating Object Storage



Pool creation and query

```
[root@client1 ~]# dm-g pool create --size=1G test
Creating DAOS pool with automatic storage allocation: 1.0 GB total, 6,94 tier ratio
Pool created with 100.00%,0.00% storage tier ratio
-----
UUID                : 71e92a47-72bd-47ab-87ce-f93bda4012f4
Service Ranks       : 0
Storage Ranks        : 0
Total Size           : 1.0 GB
Storage tier 0 (SCM) : 1.0 GB (1.0 GB / rank)
Storage tier 1 (NVM) : 0 B (0 B / rank)

[root@client1 ~]# dm-g pool query test
Pool 71e92a47-72bd-47ab-87ce-f93bda4012f4, ntarget=2, disabled=0, leader=0, version=1
Pool space info:
- Target(UOS) count:2
- Storage tier 0 (SCM):
  Total size: 1.0 GB
  Free: 1.0 GB, min:500 MB, max:500 MB, mean:500 MB
- Storage tier 1 (NVM):
  Total size: 0 B
  Free: 0 B, min:0 B, max:0 B, mean:0 B
Rebuild idle, 0 objs, 0 recs
[root@client1 ~]# _
```

Container creation

```
[root@client1 ~]# daos cont create test --label test_cont
external ERR # [7253.044249] mercury->msg: [error] /builddir/build/BUILD/me
rcury-2.1.0rc4/src/na/na_ofi.c:3047
# na_ofi_msg_send(): fi_tsend() failed, rc: -2 (No such file or directory)
external ERR # [7253.044389] mercury->hg: [error] /builddir/build/BUILD/mer
cury-2.1.0rc4/src/mercury_core.c:2727
# hg_core_forward_na(): Could not post send for input buffer (NA_NOENTRY)
hg ERR src/cart/crt_hg.c:1104 crt_hg_req_send_cb(0x1a88b70) [opc=0x102000
4 (DAOS) rpcid=0x4eae3f3b00000000 rank:tag=0:0] RPC failed; rc: DER_HG(-1020
): 'Transport layer mercury error'
mgmt ERR src/mgmt/cli_mgmt.c:882 dc_mgmt_pool_find() test: failed to get PS
replicas from 1 servers, DER_HG(-1020): 'Transport layer mercury error'
pool ERR src/pool/cli.c:198 dc_pool_choose_svc_rank() 00000000:test: dc_mgm
t_pool_find() failed, DER_HG(-1020): 'Transport layer mercury error'
pool ERR src/pool/cli.c:503 dc_pool_connect_internal() 00000000:test: canno
t find pool service: DER_HG(-1020): 'Transport layer mercury error'
ERROR: daos: DER_HG(-1020): Transport layer mercury error
[root@client1 ~]# |
```

Objective 3

Demonstrate User Interactions

Verifying Pool by Admin

```
[root@client1 ~]# dmgs pool list
Pool Size    Used Imbalance Disabled
-----
test 1.0 GB 0%    0%          0/2
```

```
[root@client1 ~]# dmgs storage query usage
Hosts          SCM-Total SCM-Free SCM-Used NVMe-Total NVMe-Free NVMe-Used
-----
--
192.168.100.8 17 GB    16 GB    7 %      0 B       0 B       N/A
```

Verifying Pool on server

- Server's storage used

```
[root@server1 ~]# df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        1.4G   0    1.4G   0% /dev
tmpfs           1.4G   0    1.4G   0% /dev/shm
tmpfs           1.4G  8.6M   1.4G   1% /run
tmpfs           1.4G   0    1.4G   0% /sys/fs/cgroup
/dev/mapper/centos-root 47G   2.5G   44G    6% /
/dev/sda1       1014M  151M   864M   15% /boot
tmpfs           278M   0    278M   0% /run/user/0
tmpfs           16G   1.2G   15G    8% /mnt/daos0
```

- Verified by Pool Id

```
[root@client1 ~]# dmg pool query test_pool
Pool 66c527c4-d7e8-437f-979d-3e790512a6b3,
Pool space info:
- Target(VOS) count:2
- Storage tier 0 (SCM):
  Total size: 1.0 GB
  Free: 1.0 GB, min:500 MB, max:500 MB, mea
- Storage tier 1 (NVMe):
  Total size: 0 B
  Free: 0 B, min:0 B, max:0 B, mean:0 B
Rebuild idle, 0 objs, 0 recs
[root@client1 ~]#
```

```
[root@server1 ~]# du -h /mnt/daos0
1.1G    /mnt/daos0/66c527c4-d7e8-437f-979d-3e790512a6b3
0       /mnt/daos0/ZOMBIES
0       /mnt/daos0/NEWBORNS
128M    /mnt/daos0/daos_sys
0       /mnt/daos0/control_raft/snapshots
32K     /mnt/daos0/control_raft
1.2G    /mnt/daos0
[root@server1 ~]#
```


Client - Client interaction

- Created test pool on Client1

```
[root@client1 ~]# dmgl pool create --size=0.2G test --ranks=0
Creating DAOS pool with automatic storage allocation: 200 MB total, 6,94 tier ratio
Pool created with 100.00%,0.00% storage tier ratio
-----
UUID                : 2cb62f35-678f-4629-8632-b65452f2af0f
Service Ranks       : 0
Storage Ranks        : 0
Total Size           : 200 MB
Storage tier 0 (SCM) : 200 MB (200 MB / rank)
Storage tier 1 (NVMe): 0 B (0 B / rank)
```

- Querying tank pool created by Client2

```
[root@client1 ~]# dmgl pool query tank
Pool ade79990-fbd6-453e-8818-f769af92122f, ntarg=2, disabled=0, leader=0, version=1
Pool space info:
- Target(VOS) count:2
- Storage tier 0 (SCM):
  Total size: 200 MB
  Free: 200 MB, min:100 MB, max:100 MB, mean:100 MB
- Storage tier 1 (NVMe):
  Total size: 0 B
  Free: 0 B, min:0 B, max:0 B, mean:0 B
Rebuild idle, 0 objs, 0 recs
```

- Created tank pool on Client2

```
[root@client2 ~]# dmgl pool create --size=0.2G tank --ranks=0
Creating DAOS pool with automatic storage allocation: 200 MB total, 6,94 tier ratio
Pool created with 100.00%,0.00% storage tier ratio
-----
UUID                : ade79990-fbd6-453e-8818-f769af92122f
Service Ranks       : 0
Storage Ranks        : 0
Total Size           : 200 MB
Storage tier 0 (SCM) : 200 MB (200 MB / rank)
Storage tier 1 (NVMe): 0 B (0 B / rank)
```

- Querying test pool created by Client1

```
[root@client2 ~]# dmgl pool query test
Pool 2cb62f35-678f-4629-8632-b65452f2af0f, ntarg=2, disabled=0, leader=0, version=1
Pool space info:
- Target(VOS) count:2
- Storage tier 0 (SCM):
  Total size: 200 MB
  Free: 200 MB, min:100 MB, max:100 MB, mean:100 MB
- Storage tier 1 (NVMe):
  Total size: 0 B
  Free: 0 B, min:0 B, max:0 B, mean:0 B
Rebuild idle, 0 objs, 0 recs
```

Both client pool verification on storage

- Verification by admin

```
[root@client2 ~]# dmg pool list
Pool Size   Used Imbalance Disabled
-----
test 200 MB 0%    0%        0/2
tank 200 MB 0%    0%        0/2
```

- Verification by server

```
[root@server ~]# du -h /mnt/daos0
8.0K    /mnt/daos0/control_raft/snapshots
52K     /mnt/daos0/control_raft
319M    /mnt/daos0/ade79990-fbd6-453e-8818-f769af92122f
319M    /mnt/daos0/2cb62f35-678f-4629-8632-b65452f2af0f
8.0K    /mnt/daos0/NEWBORNS
8.0K    /mnt/daos0/ZOMBIES
16K     /mnt/daos0/lost+found
129M    /mnt/daos0/daos_sys
766M    /mnt/daos0
```

Challenges faced

- SPDK Binding Failed:
 - Why? → Memory scan failure during DAOS server startup

```
[root@localhost ~]# systemctl status daos_server
■ daos_server.service - DAOS Server
   Loaded: loaded (/usr/lib/systemd/system/daos_server.service; disabled; vendor preset: disabled)
   Active: active (running) since Sat 2024-06-29 19:21:37 IST; 43s ago
   Main PID: 1748 (daos_server)
   CGroup: /system.slice/daos_server.service
           └─1748 /usr/bin/daos_server start

Jun 29 19:21:38 localhost.localdomain daos_server[1748]: ERROR: spdk: code = 803 description = ... "
Jun 29 19:21:38 localhost.localdomain daos_server[1748]: ERROR: mem scan failed: spdk: code = ... "
```

- Resolution: change the VM to one that supports AVX(Advanced Vector Extensions)

Challenges faced

- Server failure:
 - Resolution: Corrected the indentation in the yaml files
- Connection refused to admin:
 - Why? → the nodes had the same ip address due to NAT adapter
 - Resolution: using NAT network adapter, VM to VM connection can be established. Noted the new ip address of the server and changed the access point in all yaml files.

Challenges faced

- Client couldn't reach server:
 - Why? → Port 10001 was not listening
 - Add port 10001 in firewall configuration
- Server and agent not automatically starting on reboot: enabled the service
 - `systemctl enable daos_server`
 - `systemctl enable daos_agent`
- Changed hostnames of VMs to avoid confusion
 - `hostnamectl set-hostname new_hostname`

Challenges faced

- Formatting disappeared on reboot:
 - Why? → the file system mounted is tmpfs.
 - tmpfs (temporary file system) is a type of file system that resides in the system's RAM
 - It allows you to create and store files in memory rather than on disk.
 - Resolution: Use dcpm class
 - class: dcpm
 - scm_list: [/dev/sdb] or [/dev/md0] (for Raid 1)

Automating Server and Client Setup

Server Side

```
root@server1:~  X  root@client1:~
#!/bin/bash

# Command to reload systemd
systemctl daemon-reload

# Command to start daos_server
systemctl start daos_server

# Command to check status of daos_server
systemctl status daos_server

# Command to check disk space on /mnt/daos0
df -h /mnt/daos0
```

Client Side

```
root@server1:~ x root@client1:~ + v
#!/bin/bash

# Start DAOS Agent
echo "Starting DAOS Agent..."
systemctl start daos_agent
sleep 5 # Wait for DAOS Agent to start

# Check DAOS Agent status
echo "Checking DAOS Agent status..."
systemctl status daos_agent

# Wait a bit to ensure DAOS Agent is fully up
sleep 5

# Check DAOS Agent log
echo "DAOS Agent log:"
cat /tmp/daos_agent.log

# Query DAOS system
echo "Querying DAOS system..."
dmg system query -v

# Check if pool exists, if not, create it
POOL_NAME="test_pool"
if dmgl pool query $POOL_NAME > /dev/null 2>&1; then
    echo "Pool '$POOL_NAME' already exists."
else
    echo "Creating DAOS pool '$POOL_NAME' with automatic storage allocation..."
    dmgl pool create --size=1G $POOL_NAME
fi

# Query created pool
echo "Querying pool '$POOL_NAME'..."
dmgl pool query $POOL_NAME
```


Thank You!